

## APPENDIX

### A. Physical Concepts Details

In this section, we provide details on how we define each of our ten physical concepts, which we communicate to crowd-workers before annotation. We also list the pre-defined options provided to annotators for categorical concepts.

#### Continuous-Valued, Applicable to All Objects

**Mass:** This refers to how heavy an object is. If an object has contents inside, this includes how heavy both the object and its contents are combined.

**Fragility:** This refers to how easily an object can be broken or damaged. An object has higher *fragility* than another if a person would handle it more carefully to avoid breaking it.  
**Deformability:** This refers to how easily an object can change shape without breaking. An object has more *deformability* than another if less force is needed to change its shape without breaking it.

**Density (held-out):** This refers to the amount of mass per unit of volume of the object. If an object has contents inside, this only refers to the object, not the contents.

#### Continuous-Valued, Applicable to Containers

**Liquid Capacity (held-out):** This refers to the volume of liquid a container can contain without spilling.

#### Categorical-Valued, Applicable to All Objects

**Material:** This refers to what an object is made of. If an object is made of multiple materials, it refers to what material makes up the largest portion of the object that is visible. This does not refer to the contents of a container. The pre-defined options we include are *plastic*, *glass*, *ceramic*, *metal*, *wood*, *paper*, *fabric*, *food*, *unknown*, and *other* (annotator provides an open-ended response if this option is chosen).

**Transparency:** This refers to how much can be seen through an object. The pre-defined options we include are *transparent*, *translucent*, *opaque*, and *unknown*. *Transparent* refers to an object that can be seen clearly through, almost as if it was not there. *Translucent* refers to an object where some details can be seen through the object, but the details are not as clear as if it was *transparent*. *Opaque* refers to an object that cannot be seen through at all. This concept only refers to the object itself, and not the contents of a container. If different parts of an object have different levels of *transparency*, it refers what level applies to the largest visible portion of the object.

#### Categorical-Valued, Applicable to Containers

**Contents:** This refers to the contents of a container that are clearly visible and identifiable. The pre-defined options we include are *nothing*, *water*, *food*, *oil*, *soap*, *unknown*, and *other* (annotator provides an open-ended response if this option is chosen).

**Can Contain Liquid:** This refers to if a container can be used to transport a liquid across a room without a person

needing to be particularly careful about not spilling it. The pre-defined options we include are *yes*, *no*, and *unknown*.

**Is Sealed:** This refers to if a container can be rotated by any amount in any direction without spilling its contents. The pre-defined options we include are *yes*, *no*, and *unknown*.

**Container Categories.** We define the following categories from EgoObjects as containers: *bottle*, *container*, *plate*, *bowl*, *mug*, *water glass*, *measuring cup*, *wine glass*, *tea cup*, *frying pan*, *flowerpot*, *tin can*, *kettle*, *vase*, *coffee cup*, *mixing bowl*, *saucer*, *jug*, *serving tray*, *pitcher (container)*, and *picnic basket*.

### B. Automatic Annotation Details

We list the object categories we assign to *high* and *low* tiers for automating preference pair annotations for continuous concepts in Table VIII. We list the object categories for which we assign a concept label in Table IX. If a concept is not listed in these tables, we do not provide automatic annotations for that concept.

We originally assigned the label *metal* for *material* to the object category *house/car key*, but realized after crowdsourcing that not all instances of this category should have been given this assignment. Therefore, we manually labeled these examples for *material*, but still considered these to be automatic annotations for dataset purposes.

### C. Crowd-Sourcing Details

**Choosing Annotation Images.** There are multiple bounding box images in EgoObjects for each object instance. To determine which to present for annotating an object, we choose the bounding box with the highest CLIP [22] similarity with the object’s category label, as a heuristic for the object’s visibility. Specifically, we use the CLIP-ViT-H-14-laion2B-s32B-b79K model from OpenCLIP [39]. In Fig. 7, we show an example of a randomly sampled set of bounding boxes for an instance of category *guitar*, arranged from left-to-right in decreasing order of CLIP similarity. The objects in bounding boxes with lower CLIP similarities tend to be less visible.



Fig. 7: Bounding boxes for an instance of *guitar*, in decreasing order of CLIP similarity

**Attention Checks.** We generate attention checks for crowd-workers by randomly sampling from the automatic annotations, which have known labels. For the concepts *contents*, *density*, and *liquid capacity*, for which there are no automatic annotations, we manually label a small set of objects for attention checks.

**Other Details.** Each annotation job on Prolific consisted of 250 annotations for a single concept, of which 25 are attention checks. Participants were paid an average of 15.50

<b>Concept</b>	<b>High</b>	<b>Low</b>
Mass	television, microwave oven, table, nightstand, chest of drawers	pen, paper, spoon, fork, glasses, sunglasses, scissors, watch, necklace, house/car key, pencil, earrings, ring, screwdriver, book, container, plate, bowl, pillow, remote control, clothing, mug, laptop, knife, mobile phone, toy, computer mouse, water glass, towel, headphones, spatula, frying pan, measuring cup, banana, wallet, blanket, candle, apple, wine glass, picture frame, computer keyboard, game controller/pad, tea cup, tin can, handbag, whisk, orange, belt, plastic bag, salt and pepper shakers, cutting board, perfume, stapler, footwear, tablet coputer, teddy bear, cookie, scarf, coffee cup, ball, mixing bowl, pear, alarm clock, light switch, bread, jacket, tennis ball, sandal, saucer, laptop charger, camera, yoga mat, power plugs and sockets, cream, shirt, baseball bat, sun hat, paper towel, kitchen knife, doll, can opener, sock, facial tissue holder, boot, hair dryer
Fragility	water glass, television	house/car key, dumbbell, screwdriver, kitchen knife
Deformability	pillow, clothing, towel, blanket, belt, plastic bag, scarf, jacket, yoga mat, shirt, paper towel, sock	remote control, mug, mobile phone, computer mouse, water glass, frying pan, flowerpot, scissors, wine glass, house/car key, dumbbell, cutting board, microwave oven, toaster, blender, pressure cooker, kitchen knife, table, spoon, laptop, knife, fork, glasses, spatula, sunglasses, chair, measuring cup, pencil, picture frame, computer keyboard, game controller/pad, tea cup, tin can, salt and pepper shakers, television, cofeemaker, stapler, tablet computer, kettle, vase, coffee cup, mixing bowl, computer monitor, stool, ring, alarm clock, light switch, saucer, printer, screwdriver, guitar, camera, jug, gas stove, baseball bat, humidifier, chest of drawers, sink, can opener, nightstand, hair dryer

TABLE VIII: Object category assignments to *high* and *low* tiers for continuous concepts

Concept	Label	Categories
Material	Plastic	remote control, computer mouse, computer keyboard, game controller/pad, plastic bag
	Glass	water glass, wine glass
	Metal	tin can, kitchen knife, can opener
	Paper	book, paper, paper towel
	Fabric	clothing, towel, blanket, scarf, sock
	Food	banana, apple, orange, cookie, pear, bread
Transparency	Transparent	wine glass
	Opaque	book, pillow, remote control, clothing, laptop, mobile phone, towel, headphones, spatula, chair, frying pan, banana, wallet, flowerpot, scissors, apple, houseplant, house/car key, pencil, computer keyboard, tin can, whisk, dumbbell, orange, belt, cutting board, toaster, teddy bear, tablet computer, cookie, pear, computer monitor, stool, light switch, bread, pressure cooker, scarf, laptop charger, guitar, camera, yoga mat, shirt, baseball bat, paper towel, kitchen knife, sink, chest of drawers, can opener, boot, nightstand, hair dryer
Can Contain Liquid	Yes	bottle, mug, water glass, measuring cup, wine glass, tea cup, kettle, coffee cup, mixing bowl, jug, pitcher (container), tin can
	No	picnic basket, serving tray
Is Sealed	No	plate, bowl, mug, water glass, measuring cup, wine glass, tea cup, frying pan, flowerpot, kettle, vase, coffee cup, mixing bowl, saucer, jug, serving tray, pitcher (container), picnic basket

TABLE IX: Concept label assignments of object categories for categorical concepts

US dollars per hour, and each annotation job took on average 20-30 minutes to complete, depending on the concept.

In the annotation user interface, for each object example, the object is shown in the context of its surrounding scene, with the object indicated by its bounding box. We also provide the object’s category label to help clarify which object is to be annotated. Crowd-workers can choose an annotation label by clicking on an associated button, or typing an associated keyboard key. We also provide a *back* option to go to the previous example to correct mistakes. For the concepts *material* and *contents*, the user may choose *other* as an option, whereupon they are presented with a text box to type an open-ended label. We do not annotate objects from the categories *pet*, *cat*, and *dog*, to omit objects that are living.

We provide instructions to annotators that are specific to each concept, to encourage annotations that agree with our concept definitions. We provide an image of the instruction page provided to annotators for the *fragility* concept, which also includes an example of the annotation user interface, in Fig. 8. The instructions for how to annotate each property are also repeated at the bottom of the annotation user interface.

We detail the number of examples per concept and dataset split for PHYSOBJECTS in Table X. This is before any preprocessing of the data for annotator agreement or labels.

For the crowd-sourced data, the count refers to the number of examples, not the number of annotations, for which there are three times as many. We also provide the percent of crowd-sourced examples with majority agreement and unanimous agreement per concept in Table XI.

#### D. Training Details

**Hyperparameters.** We provide hyperparameters used for fine-tuning InstructBLIP in Table XII. These hyperparameters are largely derived from those proposed for fine-tuning BLIP-2 [10]. When fine-tuning, we only update the Q-Former parameters, as done during instruction tuning for InstructBLIP. We use a linear warmup of the learning rate, followed by a linear decay with a minimum learning rate of 0. We fine-tune using mixed precision bfloat16 training. We use a prompt template for questions, which is used both during training and inference. We load the InstructBLIP model using the LAVIS library [40]. We train and evaluate using the evaluation image processor provided by LAVIS, as we do not use image data augmentation.

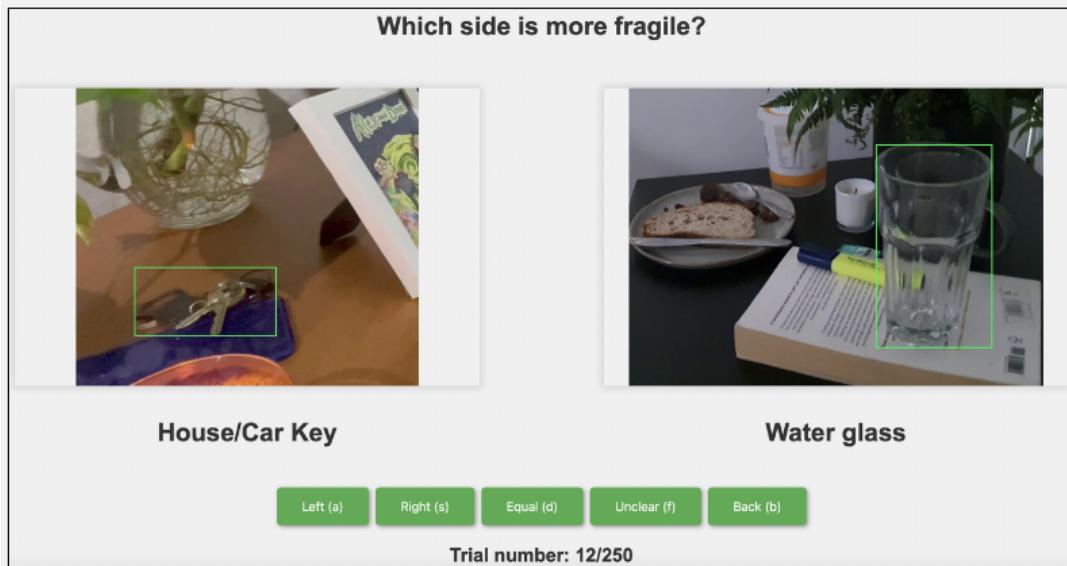
**Validation & Data Filtering.** For most experiments, we evaluate on validation data every 250 gradient steps and choose the checkpoint with the lowest validation loss. For experiments fine-tuning for a single concept, we validate every 100 gradient steps. Our validation set consists of all validation data for all concepts without balancing, except

# Physical Properties Annotations

## Instructions

Thank you for participating in this user study! In this task, you will classify 250 pairs of images based on how **fragile** the objects are.

When you begin this study, you will see an interface similar to the one below:



You will be presented with a pair of images, each of which contains an object boxed in green. There is a label below the images to help confirm what specific object we want you to analyze. **Only** use the label to identify the object, **not** to inform your annotation. If there are multiple objects of the given label in the green box, choose the object that best matches the box.

Below these are options for you to click. You may also use your keyboard to indicate an option, using the letter in parenthesis for each button. If you make a mistake, you can click the green **Back** button (or press 'b'), which allows you to go back to the previous example.

If you consider one of the objects to be more fragile than the other, you should indicate which one is more fragile using either **Left** or **Right**. If you think both objects are roughly the same fragile, choose **Equal**. If you cannot determine from the images due to low resolution, bad lighting, or other reasons, choose **Unclear**.

We define fragility as **how easily an object can be broken or damaged**. You should consider an object more fragile than another if you would handle it more carefully to avoid breaking it.

In general, please try to give some annotation instead of **Unclear** if you are somewhat confident in your responses. Even if you may be wrong, try your best to give an answer.

Occasionally, you will randomly be given an example that has a known correct answer. **You must achieve at least 80% accuracy on these examples to be compensated for the study.**

Once you complete all 250 examples, you are done, and you will be given a link to return back to Prolific to confirm your submission.

## Ready to start?

We see that your Prolific ID is **EXAMPLE**. Please contact the researchers if this is not the case.

### Consent Confirmation:

Do you consent to the collection of your responses and a temporary collection of your Prolific ID?

Yes

Fig. 8: Instruction page for the *fragility* concept

<b>Concept</b>	<b>Source</b>	<b>Train</b>	<b>Validation</b>	<b>Test</b>
Mass	Crowd-sourced	2108	86	56
	Automatic	87269	4536	2688
Fragility	Crowd-sourced	2096	99	57
	Automatic	2397	110	80
Deformability	Crowd-sourced	2101	84	65
	Automatic	293540	13384	9888
Material	Crowd-sourced	2316	460	374
	Automatic	612	130	119
Transparency	Crowd-sourced	1993	394	313
	Automatic	1046	224	194
Contents	Crowd-sourced	641	134	125
	Automatic	0	0	0
Can Contain Liquid	Crowd-sourced	318	68	64
	Automatic	342	70	67
Is Sealed	Crowd-sourced	164	30	31
	Automatic	444	91	86
Density ( <i>held-out</i> )	Crowd-sourced	0	0	500
Liquid Capacity ( <i>held-out</i> )	Crowd-sourced	0	0	500

TABLE X: Number of examples per concept and dataset split

<b>Concept</b>	<b>% Majority Agreement</b>	<b>% Unanimous Agreement</b>
Mass	94.2	58.8
Fragility	93.6	53.1
Deformability	90.5	48.1
Material	93.7	59.4
Transparency	97.0	72.5
Contents	90.4	49.8
Can Contain Liquid	99.3	64.2
Is Sealed	98.2	74.7
Density ( <i>held-out</i> )	93.3	50.7
Liquid Capacity ( <i>held-out</i> )	89.1	46.0

TABLE XI: Agreement among crowd-workers per concept

we limit the number of automatically generated examples for *mass* and *deformability* to 100. For validation data, we only use the bounding box image with the highest CLIP object category similarity score for each object, which for crowd-sourced data is also the bounding box image presented for annotation. For crowd-sourced validation data, we filter our data to only include examples with at least 2/3 majority agreement among annotators, and only use the majority label. We do not apply this filtering for training data. For preference pair annotations, we remove data annotated with *unclear*.

**Dataset Balancing.** For the categorical concepts except *is sealed*, we combine the crowd-sourced and automatically annotated data for each concept into one dataset per concept. For the other concepts, we keep separate datasets for crowd-sourced and automatically annotated data. We keep separate datasets for *is sealed* because for its crowd-sourced data, we only train using the bounding box image for the object that was presented for annotation, rather than randomly

sampling one of its bounding box images. We keep separate datasets for the continuous concepts because there is a large imbalance between the number of crowd-sourced and automatically annotated examples for these concepts. To address dataset size imbalances, we sample from each of these datasets during training at a rate proportional to the square root of the number of annotations in each dataset, as proposed in InstructBLIP for instruction tuning.

**Additional Training Details.** For most objects, each time we sample one for training, we randomly sample one of its bounding box images as input to the model, as a form of data augmentation. We do not do this with crowd-sourced data for the *contents* and *is sealed* concepts, because labels for these concepts may vary across different images of the same object. Instead, we only use the bounding box image that was presented for annotation.

To promote robustness to different queries to the VLM, we include object category labels in the question prompt for half

Hyperparameter	Value
Max fine-tuning steps	10000
Warmup steps	1000
Learning rate	1e-5
Batch size	128
AdamW $\beta$	(0.9, 0.999)
Weight decay	0.05
Image resolution	224
Prompt	Question: {} Respond unknown if you are not sure. Short answer:

TABLE XII: Hyperparameters for fine-tuning InstructBLIP

of the training examples (e.g., asking “Is this *bottle* heavy?”), and omit this information in the other half (e.g., asking “Is this *object* heavy?”). We experimented with training on one or multiple question prompts per concept, and found this to not significantly affect performance, so we only use one prompt per concept for simplicity. We include the question prompts for each concept in Table XIII. These are versions of the prompts without object category labels. When including category labels, we replace either the word “object” or “container” with the object’s category label from EgoObjects. We also pluralize the prompt to have correct grammar if the category label is plural.

We experimented with removing Q-Former text conditioning in InstructBLIP while fine-tuning, and found this to improve results on general VQA evaluation and evaluation with held-out paraphrased question prompts, so we report results using models trained without this text conditioning. In our ablation results in Table XXI, we find that this does not significantly change performance for our main crowd-sourced evaluation.

#### E. PhysObjects Evaluation

**Further Evaluation Details.** For crowd-sourced test evaluation data, we only include examples with at least 2/3 annotator agreement, and use the majority label as ground-truth. For categorical concepts, we predict by choosing the label with the highest likelihood out of all labels in PHYSOBJECTS for the concept. For continuous concepts, we predict the object in a pair with the higher score from Section IV as the one with higher concept value. We only evaluate on preference examples with a definite, non-equal preference label. For the *Most Common* baseline with continuous concepts, we also only include examples with a definite, non-equal preference when determining the most common label in the training data. We note that that *Most Common* baseline is not particularly meaningful for continuous concepts, because the preference labels and predictions are invariant to ordering in each preference pair. Therefore, a more natural baseline for these concepts would be random guessing, which would achieve 50% accuracy.

Similarly as with validation data, for test data we only evaluate using the bounding box image with the highest CLIP object category similarity per object, which for crowd-sourced data is also the bounding box image presented for

annotation. We evaluate using the same question prompts per concept as during training, which are listed in Table XIII. We generally report evaluation results without object category labels in the question prompt, because this gives slightly better results for the base InstructBLIP model.

**Text Only Baseline.** For this baseline, we use ground truth object category labels from EgoObjects. We use the ‘text-davinci-003’ InstructGPT model [41] as our LLM. For each concept, we use 128 in-context examples randomly sampled from the training data in PHYSOBJECTS for that concept. Because in-context learning is limited by context length, and therefore it is desirable to use the best quality in-context examples when possible, we first apply to the training data the same majority filtering process used on crowd-sourced test data as described in the previous subsection. We also remove preference annotations with the label *unclear*, as done in our VLM fine-tuning setup. We treat each example as a question answering task, using question prompts for each concept similar to those in Table XIII, but modified to refer to general object classes, rather than specific instances. We make predictions by selecting the most likely completion of the LLM conditioned on the in-context examples and test example. For categorical concepts, we first include in the LLM context all possible labels in PHYSOBJECTS for the concept. For continuous concepts, because we only evaluate on examples with definite preferences, we restrict predictions to only definite preferences using logit bias, although the in-context examples include *equal* as a possible answer.

**Paraphrased Question Prompts.** In Table XIV, we list the paraphrased prompts used in the evaluation for Table IV.

**Limited VQA Degradation.** Ideally, training on PHYSOBJECTS should be done while co-training on other vision and language datasets to preserve general reasoning abilities. In this work, we do not do this because we focus primarily on physical reasoning. However, we show that fine-tuning on only PHYSOBJECTS does not significantly degrade general VQA performance. In Table XV, we compare InstructBLIP to PG-InstructBLIP on VQAv2 [42] and OK-VQA [43]. These results suggest that existing systems using VLMs can benefit from PHYSOBJECTS for physical reasoning, without sacrificing other reasoning abilities.

We perform VQA evaluation using the LAVIS library, using their configurations for evaluation of BLIP-2. Although PG-InstructBLIP is fine-tuned without Q-Former text condi-

Concept	Question Prompt
Mass	Is this object heavy?
Fragility	Is this object fragile?
Deformability	Is this object deformable?
Material	What material is this object made of?
Transparency	Is this object transparent, translucent, or opaque?
Contents	What is inside this container?
Can Contain Liquid	Can this container hold a liquid inside easily?
Is Sealed	Is this container sealed?
Density ( <i>held-out</i> )	Is this object dense?
Liquid Capacity ( <i>held-out</i> )	Can this object hold a lot of liquid?

TABLE XIII: Question prompts for each concept, without object category labels

Concept	Question Prompt
Mass	Does this object weigh a lot?
Fragility	Is this object easily breakable?
Deformability	Is this object easily bendable?
Material	What is this object made of?
Transparency	Would you describe this object as opaque, transparent, or translucent?
Contents	What does this container contain?
Can Contain Liquid	Is this container able to hold water inside easily?
Is Sealed	Is this container sealed shut?

TABLE XIV: Paraphrased question prompts for main concepts, without object category labels

	InstructBLIP	PG-InstructBLIP (ours)
VQAv2	71.4	67.5
OK-VQA	52.4	48.7

TABLE XV: Accuracy on existing VQA benchmarks

tioning, we found that Q-Former text conditioning during VQA evaluation improved performance, so we report these results. We believe this is because InstructBLIP was instruction tuned with this text conditioning. We also experimented with VQA evaluation on PG-InstructBLIP fine-tuned with Q-Former text conditioning, but found this to have worse results, likely due to overfitting on our limited variety of question prompts. We believe these issues can be mitigated by co-training on PHYSOBJECTS in combination with other vision and language datasets, which we leave for future work.

Motivated by these VQA results, for our planning evaluations we also evaluate PG-InstructBLIP using Q-Former text conditioning, to avoid possible degradation when answering questions that do not pertain concepts in PHYSOBJECTS. We also verified that evaluating PG-InstructBLIP using Q-Former text conditioning did not significantly affect test accuracy on PHYSOBJECTS.

**Including Object Category Labels in Question Prompts.** We generally report evaluation results without ground-truth object category labels in the question prompt. In Table XVI, we compare including object category labels or not, and find that all models are not extremely sensitive to this.

**Including Concept Definitions in Question Prompts.**

While we did not spend extensive effort designing the question prompts for each concept (shown in Table XIII), we aimed for them to be concise while still eliciting the desired concept. As seen in Table XVIII, the base InstructBLIP model achieves above chance performance on all concepts, suggesting that these prompts do elicit the desired concept to some extent. However, these prompts do not contain our definitions for each concept provided to annotators, as described in Appendix A. We analyze whether including concept definitions in the question prompt would improve base VLM performance in Table XVIII, which contains our original crowd-sourced test accuracy results, with additional evaluation of the base InstructBLIP model using modified prompts that contain concept definitions, which we provide in Table XVII. We find that while including concept definitions improves performance for some concepts (*mass*, *deformability*, *contents*, *can contain liquid*), this still does not match PG-InstructBLIP on these concepts, and overall performance in fact *decreases* compared to the original prompts. We believe this could be because InstructBLIP does not have strong enough language understanding to properly incorporate the concept definitions when providing responses. For this reason, and for simplicity, we use the original prompts without concept definitions in our experiments, unless otherwise stated.

**Using a Smaller VLM.** To analyze the effect of VLM size on physical reasoning, in Table XIX we provide evaluation results using the InstructBLIP version with the smaller Flan-T5 XL as its base LLM, compared to the Flan-T5 XXL

	InstructBLIP		Single Concept FT (ours)		PG-InstructBLIP (ours)	
Category Labels	Yes	No	Yes	No	Yes	No
Mass	60.0	62.2	<b>84.4</b>	80.0	80.0	80.0
Fragility	75.7	78.4	91.2	91.2	<b>97.3</b>	94.6
Deformability	69.8	67.4	88.4	<b>95.3</b>	90.7	93.0
Material	73.3	67.1	<b>86.8</b>	83.7	85.7	84.6
Transparency	84.5	85.8	89.1	89.4	89.8	<b>90.1</b>
Contents	34.2	35.1	80.7	81.6	82.5	<b>83.3</b>
Can Contain Liquid	57.8	59.4	84.4	84.4	82.8	<b>87.5</b>
Is Sealed	71.0	74.2	80.6	80.6	<b>87.1</b>	<b>87.1</b>
Average	65.8	66.2	85.7	85.8	87.0	<b>87.5</b>

TABLE XVI: Test accuracy for main concepts on crowd-sourced PHYSOBJECTS, with and without object category labels

Concept	Question Prompt
Mass	The heaviness of an object refers to its mass. It includes the contents of the object if it has something inside it. Is this object heavy?
Fragility	Fragility refers to how easily an object can be broken or damaged. Is this object fragile?
Deformability	Deformability refers to how easily an object can change shape without breaking. Is this object deformable?
Material	The material of an object refers to what material makes up the largest portion of the object that is visible. It does not refer to the contents of a container. What material is this object made of?
Transparency	Transparency refers to how much can be seen through an object. A transparent object can be clearly seen through, almost as if it was not there. A translucent object can be seen through with some details, but not as clearly as if it was transparent. An opaque object cannot be seen through at all. The transparency of an object does not refer to the transparency of its contents if it has anything inside it. Is this object transparent, translucent, or opaque? If different portions of the object have different levels of transparency, respond with the level that applies to the largest visible portion of the object.
Contents	What is inside this container? Only respond with contents that are clearly visible and identifiable.
Can Contain Liquid	A container can contain liquid if it can be used to transport a liquid across a room without a person needing to be particularly careful about not spilling it. Can this container contain liquid?
Is Sealed	A container is sealed if it can be rotated by any amount in any direction without spilling its contents if it has anything inside. Is this container sealed?

TABLE XVII: Question prompts with definitions for each main concept, without object category labels

Prompt Type	InstructBLIP		PG-InstructBLIP (ours)
	Original	w/ Concept Definitions	Original
Mass	62.2	71.1	<b>80.0</b>
Fragility	78.4	67.6	<b>94.6</b>
Deformability	67.4	69.8	<b>93.0</b>
Material	67.1	66.0	<b>84.6</b>
Transparency	85.8	65.3	<b>90.1</b>
Contents	35.1	36.0	<b>83.3</b>
Can Contain Liquid	59.4	64.1	<b>87.5</b>
Is Sealed	74.2	64.5	<b>87.1</b>
Average	66.2	63.0	<b>87.5</b>

TABLE XVIII: Test accuracy for main concepts on crowd-sourced PHYSOBJECTS, with additional base InstructBLIP evaluation on prompts with definitions for each concept

version used in all other experiments. We find that while the smaller Flan-T5 XL version generally has worse base performance, after fine-tuning on PHYSOBJECTS, we see that performance is comparable between the two model sizes. This suggests that for physical reasoning, fine-tuning on human data such as PHYSOBJECTS could reduce the need for larger model sizes. While fine-tuned evaluation performance is similar across model sizes, for simplicity of comparison, we only report results using the larger Flan-T5 XXL models in all other experiments.

**Results on Automatically Annotated Data.** We report evaluation results on automatically annotated data in Table XX. Performance is generally much higher on this data compared to the crowd-sourced data, because these are easier examples that can be determined from object categories alone.

**Held-Out Concept Scaling.** To evaluate the transfer abilities of InstructBLIP across different concepts, we evaluate how fine-tuning it on PHYSOBJECTS improves performance scaling on held-out concepts. We fine-tune models on data from all concepts except one, and then report results of additional fine-tuning on the held-out concept. We compare to fine-tuning base InstructBLIP without training on the other concepts, and base InstructBLIP without fine-tuning. Results for three concepts are shown in Fig. 9. We chose these concepts because we believed they had the most generalization potential from the other concepts. We find that there are some signs of positive transfer on *mass* and *fragility*, although we see slight negative transfer on *material*. We believe that more positive transfer could be attained by co-training with other vision and language datasets.

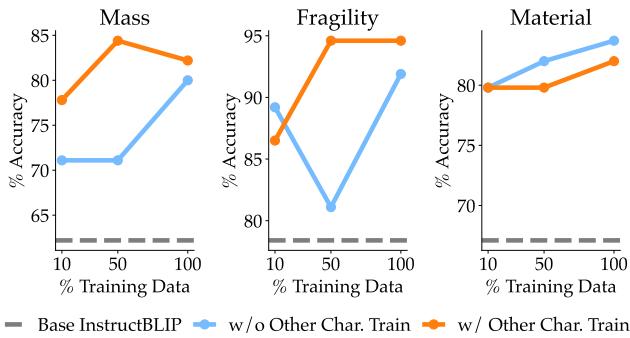


Fig. 9: Performance scaling on held-out concepts

**Ablations.** We report additional ablation results on crowd-sourced PHYSOBJECTS examples in Table XXI. We list each ablation below:

- 1) *No Auto Data*: Instead of training on both crowd-sourced and automatically annotated data, we train on only crowd-sourced data.
- 2) *Filtered*: Instead of training on all annotations for crowd-sourced data, we filter the data similarly as during evaluation: we only include examples with at least 2/3 annotator agreement, and use the majority label as ground-truth.
- 3) *Q-Former Text*: Instead of removing Q-Former text

conditioning during fine-tuning, we include it, as done for the original InstructBLIP model.

- 4) *No Category Info*: Instead of training on both question prompts with and without object category information, we only train on question prompts *without* it.
- 5) *Only Category Info*: Instead of training on both question prompts with and without object category information, we only train on question prompts *with* it. Here, unlike the rest of the evaluations, we evaluate with object category information to match the training setup.

We find that overall performance for each ablated version of our model does not change significantly, suggesting some robustness of our fine-tuning process to different design decisions. In particular, we find that including automatically annotated data does not significantly impact performance on crowd-sourced data, which perhaps is not surprising because the object categories differ. *Only Category Info* very slightly improves upon PG-InstructBLIP, but uses privileged object category information at evaluation time.

#### F. Real Scene Planning Evaluation Details

**Planning Framework.** Our planning framework consists of first providing the scene image to an OWL-ViT ViT-L/14 open-vocabulary object detector [44], which produces object bounding boxes and category labels from the EgoObjects categories. We then provide the list of detected objects and the task instruction to our LLM, which is GPT-4 [38] with temperature 0. The LLM is additionally provided with the robotic primitives, and a few-shot chain-of-thought prompt [45] with instructions to ask questions about objects in the scene to determine how to complete the task, and then produce a plan using the primitives. There is no constraint on the questions that the LLM can ask, except for encouragement in the prompt to ask questions that can be answered with *yes* or *no*. The same prompt is used for all scenes and tasks, which we provide in Listing 1.

After the LLM asks a set of object-centric questions, a VLM answers each question prompted with the bounding box of the object indicated by the LLM, and then provides the LLM with its highest likelihood responses and their associated likelihoods/confidence scores, as done in prior work for VQA [30]. This continues until the LLM decides it has enough information, whereupon it either indicates that the task is not possible, or produces a plan consisting of a list of primitives to execute for the task. The few-shot examples in Listing 1 illustrate how interaction between the LLM and VLM for planning is structured.

**Primitives.** We list the primitives for our real scene planning evaluation below:

- go to object [X]
- pick up object [X]
- bring to human object [X]
- put down object [X]
- done

The primitives (except *done*) are parameterized by a letter (in place of [X]) that identifies each detected object in the

LLM Version	InstructBLIP		PG-InstructBLIP (ours)	
	Flan-T5 XL	Flan-T5 XXL	Flan-T5 XL	Flan-T5 XXL
Mass	62.2	62.2	<b>82.2</b>	80.0
Fragility	64.9	78.4	<b>97.3</b>	94.6
Deformability	48.8	67.4	<b>97.7</b>	93.0
Material	69.1	67.1	82.6	<b>84.6</b>
Transparency	74.0	85.8	87.5	<b>90.1</b>
Contents	18.4	35.1	<b>86.8</b>	83.3
Can Contain Liquid	68.8	59.4	<b>87.5</b>	<b>87.5</b>
Is Sealed	67.7	74.2	80.6	<b>87.1</b>
Average	59.2	66.2	<b>87.8</b>	87.5

TABLE XIX: Test accuracy for main concepts on crowd-sourced PHYSOBJECTS, using different VLM versions

	InstructBLIP	PG-InstructBLIP (ours)
Mass	72.8	<b>99.9</b>
Fragility	95.0	<b>100</b>
Deformability	96.0	<b>98.8</b>
Material	89.1	<b>98.3</b>
Transparency	97.4	<b>100</b>
Can Contain Liquid	98.5	<b>100</b>
Is Sealed	<b>100</b>	<b>100</b>
Average	92.7	<b>99.6</b>

TABLE XX: Test accuracy for main concepts on automatically annotated PHYSOBJECTS

scene. The assignment of letters is provided in the list of object detections given to the LLM planner.

**Scenes and Tasks.** In Table XXII, we provide the scene images in our evaluation, and the detected objects and task instructions for each scene. We also indicate the task type for each instruction.

**Prompts.** We provide the prompts used by our LLM-based planning framework for our scene planning evaluation. The version with VLM interaction is in Listing 1 and the version without VLM interaction is in Listing 2. The parts of the prompts in square brackets are replaced with the corresponding information specific to the task, in the same format as the prompt example.

#### Listing 1: Prompt for LLM planner with VLM interaction.

You are a household robot. You are able to move most household objects, but not large or heavy furniture.  
You are to be safe and not break anything.  
You will be given a list of objects in the scene.  
A human will give you a task instruction to perform.  
First, ask questions about the objects to learn more about them to determine how to properly complete the task.  
Indicate the object letters before asking the question, and do not provide the answer.  
Format each question like "Question about object [A, B]: Is this object heavy?".  
Only ask informative questions that help you understand how to properly complete the task.  
Some of the object descriptions in the list may be inaccurate, so it is good to ask questions to confirm information about them.  
Ask questions that can be answered with yes or no when possible.  
Only ask one question at a time.

After each question, you will be given a response containing the 3 most likely possible answers for each object, with corresponding confidence scores. An example response for one object would be "A: yes (0.8), no (0.1), unknown (0.1)". Otherwise, use the information you have obtained to produce a numbered list of actions to perform the task, starting with "Plan:". You are only able to perform the following actions, where "X" is a placeholder for an object letter from the given list of objects:  
1. go to object X  
2. pick up object X  
3. bring to human object X  
4. put down object X  
5. done  
  
Only perform actions that are possible. You are able to move most household objects, including most electronic appliances, but not large or heavy furniture or parts of a room (like a door or heater). You are to be safe and not break anything. Choose the "done" action when the plan is complete and then stop. If no plan can satisfy the task, tell the human, choose the "done" action and stop.  
Only respond with the plan list and do not include any other text.  
  
Scene 1:  
The following objects are in the scene: A (bottle), B (bowl), C (countertop), D (shirt), E (lock)  
  
Instruction: Bring me the heaviest object.  
  
Thought: I can ask which object is heavy out of all the objects, and I can find the one with the largest weight that I can carry.  
Question about [A, B, C, D, E]: Is this object heavy?  
Answer:  
A: No (0.50), Yes (0.24), Unknown (0.21)  
B: No (0.90), Unknown (0.05), Yes (0.04)  
C: Yes (0.80), Unknown (0.10), No (0.05)  
D: No (0.73), Unknown (0.20), Yes (0.06)  
E: No (0.41), Yes (0.35), Unknown (0.19)  
Thought: The heaviest object is the one with the highest Yes value or the lowest No value (if none are Yes). The heaviest object is C. However, I cannot carry a countertop since it is a heavy piece of furniture. Therefore, I will choose the next heaviest, which would be E, a lock that I can carry easily.  
Plan:  
1. Go to object E  
2. Pick up object E  
3. Bring to human object E  
4. Done  
  
Instruction: Bring me the most deformable container.  
Thought: I can first ask which objects are containers out of all the objects I can carry. Of these containers, I can ask which is deformable.

	PG-InstructBLIP (ours)	No Auto Data	Filtered	Q-Former Text	No Category Info	Only Category Info
Mass	80.0	<b>84.4</b>	75.6	80.0	75.6	77.8
Fragility	94.6	94.6	97.3	97.3	<b>100</b>	<b>100</b>
Deformability	<b>93.0</b>	<b>93.0</b>	90.7	90.7	<b>93.0</b>	88.4
Material	84.6	83.4	83.4	85.4	84.6	<b>86.5</b>
Transparency	90.1	89.4	89.1	<b>92.1</b>	89.8	91.7
Contents	83.3	81.6	85.1	<b>87.7</b>	84.2	86.8
Can Contain Liquid	87.5	89.1	85.9	84.4	<b>90.6</b>	84.4
Is Sealed	<b>87.1</b>	83.9	83.9	71.0	80.6	<b>87.1</b>
Average	87.5	87.4	86.4	86.1	87.3	<b>87.8</b>

TABLE XXI: Ablation results for main concepts on crowd-sourced PHYSOBJECTS

Question about [A, B, C, D, E]: Is this object a container ?  
 Answer:  
 A: Yes (0.55), No (0.20), Unknown (0.19)  
 B: Yes (0.90), Unknown (0.05), No (0.04)  
 C: No (0.85), Unknown (0.06), Yes (0.05)  
 D: No (0.62), Unknown (0.20), Yes (0.06)  
 E: No (0.41), Yes (0.35), Unknown (0.19)  
 Thought: The only objects that are confidently Yes are B and A. I should ask which is deformable next.  
 Question about [A, B]: Is this object deformable?  
 Answer:  
 A: Yes (0.80), Unknown (0.15), No (0.04)  
 B: No (0.55), Unknown (0.26), Yes (0.05)  
 Thought: The most deformable object is the one with the highest Yes value or the lowest No value (if none are Yes). Since some of the answers are Yes, the answer is A (the highest yes value).  
 Plan:  
 1. Go to object A  
 2. Pick up object A  
 3. Bring to human object A  
 4. Done  
 Scene 2:  
 The following objects are in the scene: [list of objects in the scene]  
 Instruction: [instruction specified here]

### Listing 2: Prompt for LLM planner without VLM interaction.

You are a household robot. You are able to move most household objects, but not large or heavy furniture. You are to be safe and not break anything. You will be given a list of objects in the scene. A human will give you a task instruction to perform. Use the object information to produce a numbered list of actions to perform the task, starting with "Plan:". You are only able to perform the following actions, where "X" is a placeholder for an object letter from the given list of objects:

1. go to object X
2. pick up object X
3. bring to human object X
4. put down object X
5. done

Only perform actions that are possible. You are able to move most household objects, including most electronic appliances, but not large or heavy furniture or parts of a room (like a door or heater). You are to be safe and not break anything. Choose the "done" action when the plan is complete and then stop. If no plan can satisfy the task, tell the human, choose the "done" action and stop.

Only respond with the plan list and do not include any other text.

Scene 1:  
 The following objects are in the scene: A (bottle), B (bowl), C (countertop), D (shirt), E (lock)

Instruction: Bring me the heaviest object.  
 Thought: I cannot carry a countertop since it is a heavy piece of furniture. Out of the rest, a good guess would be the lock.  
 Plan:  
 1. Go to object E  
 2. Pick up object E  
 3. Bring to human object E  
 4. Done  
 Instruction: Bring me the most deformable container.  
 Thought: Typically shirts are easy to fold, so a good choice for the most deformable object would be the shirt.  
 Plan:  
 1. Go to object D  
 2. Pick up object D  
 3. Bring to human object D  
 4. Done  
 Scene 2:  
 The following objects are in the scene: [list of objects in the scene]  
 Instruction: [instruction specified here]

**Evaluation Procedure.** We evaluate task planning accuracy using a non-author human evaluator. For each evaluation, the evaluator is given the task instruction, the image of the scene, the list of detected objects in the scene and their bounding boxes, and the generated task plan, and they are asked to evaluate whether the task plan successfully performed the task instruction for the given scene. We provide the following instructions to the evaluator on what to consider when evaluating whether a plan was correct:

Instructions: For each scene, there is a list of objects (under 'Options:'). Below that is a table of tasks for that scene. The instruction given to a robot is on the left. On the right are the choices from 3 different robots. You need to mark which ones are correct or incorrect. It may be possible that multiple robots got it right or none of them got it right. Be aware that in tasks that involve moving objects, the robot should not plan to move an object that is very heavy, like large furniture.

While the planner usually creates plans using only the provided primitives, it sometimes specifies primitives that were not provided. Because the purpose of this evaluation is on assessing if a LLM planner can benefit from physical reasoning using a VLM, and not on creating a functional planning system, we do not do anything to handle these cases. We the evaluator to judge if these plans satisfy the task

instruction like the others. We provide example executions for different versions of our planning framework on our [website](#).

#### G. Real Robot Evaluation Details

**Primitives.** We list the primitives for our real robot evaluation below:

- move [X] to the side
- move [X] into [Y]
- done

Similarly as with our planning-only evaluation, our primitives are parameterized by a letter (in place of [X] or [Y]) that identifies each detected object in the scene. The assignment of letters is provided in the list of object detections given to the LLM planner.

**Scenes and Tasks.** In [Table XXIII](#), we provide the scene images in our real robot evaluation, the detected objects in each scene, and the task instructions for each scene.

**Prompts.** We use the same prompt structure from the real scene planning evaluation. For tasks that involve moving an object to a side, we add “In your plan, you may only use the following primitive: move X to the side (where X is an object). Do not move furniture.” For tasks that involve moving an object into a container, we add “In your plan, you may only use the following primitive: move X into Y (where X is an object and Y is a container). Do not move furniture.”

**Evaluation Procedure.** We run real robot experiments using a 7-DoF Franka Emika Panda robot with a Robotiq 2F-85 gripper, using Polymetis [46] for real-time control. We obtain our pick-and-place primitives by collecting a kinesthetic demonstration for each primitive, and replaying the demonstration to execute it. The objects in the images provided to the object detector are not in the exact same positions as when the robot is acting, because the objects have to be rearranged when collecting demonstrations for each primitive. However, this does not affect the planner because we do not provide it object positions, as our planning framework does not make use of them.

Because our evaluation focuses on planning quality and not successful execution of primitives, we retry execution of each plan until all of the primitives are successfully executed. Therefore, our success rates are only reflective of planning quality, and not that of the primitives.

We evaluate the success rate of robot executions using a non-author human evaluator. For each evaluation, the evaluator is given the task instruction and a video of the robot executing the generated plan, and they are asked to evaluate whether the robot successfully performed the task instruction. We provide visualizations of robot executions on our [website](#).

Scene Image	Object Detections	Task Instructions
	<ul style="list-style-type: none"> <li>1) bottle</li> <li>2) pitcher (container)</li> <li>3) bowl [flatter bowl]</li> <li>4) towel [shirt]</li> <li>5) countertop</li> <li>6) bowl [taller ceramic bowl]</li> <li>7) measuring cup [lock]</li> </ul>	<ul style="list-style-type: none"> <li>1) Bring me the heaviest object. [S]</li> <li>2) Bring me the most deformable object. [S]</li> <li>3) Bring me the most fragile object. [S]</li> <li>4) Bring me all containers that you can confidently determine have water. [M]</li> <li>5) Bring me the container with oil. [M]</li> <li>6) Among all empty containers, bring me the ones that cannot be used to carry water. [M]</li> <li>7) Bring me the metal object. [S]</li> </ul>
	<ul style="list-style-type: none"> <li>1) suitcase [blue crate]</li> <li>2) stool</li> <li>3) hair dryer [mirror]</li> <li>4) chair [chair that the mirror is on]</li> <li>5) dishwasher [metal cabinet in top right]</li> <li>6) chair [blue chair]</li> <li>7) bottle [Elmer glue container]</li> <li>8) bottle [Mod Podge container]</li> <li>9) container [paint thinner container]</li> <li>10) desk</li> <li>11) mug [mug with paintbrushes]</li> <li>12) facial tissue holder [container with glitter]</li> <li>13) pencil</li> </ul>	<ul style="list-style-type: none"> <li>1) Bring me the heaviest object. [S]</li> <li>2) Bring me a metal container. [M]</li> <li>3) Bring me a small, empty cup that I can fill with water to clean my paintbrushes. If there are none, tell me that there are no small empty cups. [M]</li> <li>4) Bring me the clear container with art supplies. [C]</li> <li>5) Bring me the metal object that is reflective. [M]</li> <li>6) Bring me paint thinner. [C]</li> <li>7) Bring me a wooden object. [S]</li> </ul>

TABLE XXII: Scene images, object detections, and task instructions for our real scene planning evaluation (scenes 1 and 2). The object category labels given by OWL-ViT are sometimes inaccurate or ambiguous, in which case we provide more precise labels in square brackets. Note that the planner only has access to the original OWL-ViT labels. Tasks are labeled with S, M, or C for *Single Concept*, *Multi-Concept*, or *Common Knowledge*, respectively.

Scene Image	Object Detections	Task Instructions
	<ol style="list-style-type: none"> <li>1) clothing [green hoodie]</li> <li>2) towel</li> <li>3) clothing [striped shirt]</li> <li>4) bottle [sunscreen bottle]</li> <li>5) towel [socks]</li> <li>6) mouse [ear thermometer]</li> <li>7) suitcase</li> <li>8) bottle [hand sanitizer]</li> <li>9) hair dryer [dumbbell]</li> <li>10) clothing [blue shirt]</li> </ol>	<ol style="list-style-type: none"> <li>1) Bring me the heaviest object. [S]</li> <li>2) Bring me all clear containers. [M]</li> <li>3) Bring me the hard plastic object. [M]</li> <li>4) Bring me the lightest piece of clothing. [S]</li> <li>5) Bring me the object I can pack my clothes into. [C]</li> <li>6) It is cold outside. Bring me something that can keep me warm. [C]</li> <li>7) It is sunny outside. Bring me the container of sunscreen. [C]</li> </ol>
	<ol style="list-style-type: none"> <li>1) facial tissue holder [paper towel dispenser]</li> <li>2) light switch [left electric outlet]</li> <li>3) light switch [right electric outlet]</li> <li>4) mixer</li> <li>5) toaster</li> <li>6) kettle</li> <li>7) paper towel</li> <li>8) water glass [plastic cup]</li> <li>9) salt and pepper shakers [salt]</li> <li>10) bottle [jam container]</li> <li>11) frying pan [baking pan]</li> <li>12) container [salmon-colored container]</li> <li>13) salt and pepper shakers [pepper]</li> <li>14) countertop</li> </ol>	<ol style="list-style-type: none"> <li>1) Bring me the heaviest object. [S]</li> <li>2) Bring me the heavier glass container. [M]</li> <li>3) Bring me something that is easy to tear. [C]</li> <li>4) Bring me the lightest container that is empty but can be filled with water. [M]</li> <li>5) Bring me the most deformable container with a lid. [M]</li> <li>6) Bring me all metal containers that can be used to carry water. [M]</li> <li>7) Bring me the object that can be used in an oven. [C]</li> </ol>

TABLE XXII: Scene images, object detections, and task instructions for our real scene planning evaluation (scenes 3 and 4). The object category labels given by OWL-ViT are sometimes inaccurate or ambiguous, in which case we provide more precise labels in square brackets. Note that the planner only has access to the original OWL-ViT labels. Tasks are labeled with S, M, or C for *Single Concept*, *Multi-Concept*, or *Common Knowledge*, respectively.

Scene Image	Object Detections	Task Instructions
	<ul style="list-style-type: none"> <li>1) toaster</li> <li>2) light switch [electric outlet]</li> <li>3) envelope [napkin on microwave]</li> <li>4) light switch</li> <li>5) microwave oven [microwave]</li> <li>6) door [microwave door]</li> <li>7) bottle [glass sauce bottle]</li> <li>8) picnic basket [drying rack]</li> <li>9) soap dispenser</li> <li>10) bottle [plastic bottle with blue vanilla flavor]</li> <li>11) mug [dry mug]</li> <li>12) sink</li> <li>13) frying pan [dirty pan in sink]</li> <li>14) mug [dirty mug in sink]</li> <li>15) countertop</li> <li>16) waste container</li> <li>17) cupboard</li> <li>18) plastic bag [trashbag]</li> </ul>	<ul style="list-style-type: none"> <li>1) Bring me the heaviest object that you can carry. [S]</li> <li>2) Bring me an empty mug that I can use to make tea. [C]</li> <li>3) Bring me the most deformable object. [S]</li> <li>4) Bring me the glass object. [S]</li> <li>5) Bring me a metal pan that is in the sink. [C]</li> <li>6) Bring me the container that stores trash. [C]</li> </ul>
	<ul style="list-style-type: none"> <li>1) envelope [sign on napkin dispenser]</li> <li>2) humidifier [napkin dispenser]</li> <li>3) ladle [metal tongs]</li> <li>4) food [two salad containers on the right]</li> <li>5) bottle [red wine vinegar bottle]</li> <li>6) frying pan [closer salad tray]</li> <li>7) paper [napkin coming out of dispenser]</li> <li>8) countertop</li> <li>9) bottle [olive oil bottle]</li> <li>10) bottle [black container on the right]</li> <li>11) bottle [black container on the left]</li> <li>12) juice [olive oil inside bottle]</li> <li>13) cabinetry</li> <li>14) countertop [more cropped in view of countertop]</li> <li>15) bowl [paper plate under the counter]</li> </ul>	<ul style="list-style-type: none"> <li>1) Bring me the most deformable object. [S]</li> <li>2) Bring me the lightest metal object. [M]</li> <li>3) Bring me the heaviest glass container. [M]</li> <li>4) Serve some food on a plate using objects in the scene. [C]</li> <li>5) Bring me an empty container that you can confidently use to contain liquids, if one exists. Otherwise, tell me that no suitable containers exist. [M]</li> <li>6) Bring me the container of olive oil. [C]</li> </ul>

TABLE XXII: Scene images, object detections, and task instructions for our real scene planning evaluation (scene 6). The object category labels given by OWL-ViT are sometimes inaccurate or ambiguous, in which case we provide more precise labels in square brackets. Note that the planner only has access to the original OWL-ViT labels. Tasks are labeled with S, M, or C for *Single Concept*, *Multi-Concept*, or *Common Knowledge*, respectively.

Scene Image	Object Detections	Task Instructions
	<ol style="list-style-type: none"> <li>1) whiteboard</li> <li>2) door [leftmost door]</li> <li>3) paper</li> <li>4) window [window of left door of rightmost pair]</li> <li>5) door [left door of rightmost pair]</li> <li>6) table [taller table]</li> <li>7) chair [leftmost short chair facing towards the camera]</li> <li>8) chair [tall chair]</li> <li>9) chair [short chair behind pillar]</li> <li>10) chair [rightmost short chair, facing towards the camera]</li> <li>11) table [long wooden table]</li> <li>12) door [rightmost door]</li> <li>13) couch</li> <li>14) chair [left side, facing away from camera]</li> <li>15) coffee table</li> </ol>	<ol style="list-style-type: none"> <li>1) Go to the piece of furniture that is the softest. [C]</li> <li>2) Go to the glass object that is not part of a window or door. [S]</li> <li>3) Bring me the lightest object. [S]</li> <li>4) Among all pieces of furniture, go to the one that is lightest. [C]</li> <li>5) Go to the table that does not have a wooden surface. [S]</li> </ol>
	<ol style="list-style-type: none"> <li>1) box [binder]</li> <li>2) bottle [large plastic tub]</li> <li>3) bottle [plastic bottle]</li> <li>4) box [algorithms textbook]</li> <li>5) pitcher (container) [blue metal cup]</li> <li>6) water glass [small glass cup]</li> <li>7) headphones [phone cable]</li> <li>8) dumbbell [power brick]</li> <li>9) adhesive tape [ruler]</li> </ol>	<ol style="list-style-type: none"> <li>1) Bring me the container that is most likely to be metal. [M]</li> <li>2) Bring me the heaviest container. Only consider the container itself, not the contents inside. [S]</li> <li>3) Bring me the sealed container with juice. [M]</li> <li>4) Bring me the most bendable object. [S]</li> <li>5) Bring me the most fragile object. [S]</li> <li>6) Bring me all containers that are made of plastic (with very high confidence). [M]</li> </ol>

TABLE XXII: Scene images, object detections, and task instructions for our real scene planning evaluation (scenes 7 and 8). The object category labels given by OWL-ViT are sometimes inaccurate or ambiguous, in which case we provide more precise labels in square brackets. Note that the planner only has access to the original OWL-ViT labels. Tasks are labeled with S, M, or C for *Single Concept*, *Multi-Concept*, or *Common Knowledge*, respectively.

Scene Image	Object Detections	Task Instructions
	1) towel [handbag] 2) bottle [paint bottle] 3) bowl [yellow plastic bowl] 4) tool [container of metals] 5) desk [full table] 6) saucer 7) bowl [ceramic bowl] 8) spoon 9) pencil [pen] 10) pencil 11) milk [snack packet]	1) Move all objects that are not plastic to the side. 2) Find a container that has metals. Move all metal objects into that container. 3) Move all containers that can be used to carry water to the side. 4) Put the two objects with the least mass into the least deformable container. 5) Move the most fragile object to the side.
	1) bottle [glass jar] 2) mug [metal mug] 3) scale [cardboard cupholder] 4) mug [plastic cup] 5) adhesive tape 6) tool [wrench]	1) Put all containers that can hold water to the side. 2) Put all objects that are not plastic to the side. 3) Put all objects that are translucent to the side. 4) Put the three heaviest objects to the side. 5) Put a plastic object that is not a container into a plastic container. Choose the container that you are most certain is plastic.

TABLE XXIII: Scene images, object detections, and task instructions for our real robot evaluation. The object category labels given by OWL-ViT are sometimes inaccurate or ambiguous, in which case we provide more precise labels in square brackets. Note that the planner only has access to the original OWL-ViT labels.